

Ethan Gates

(507) 316-2063 | ethangates5329@gmail.com | Rochester, MN | [linkedin.com/in/ethan-gates5329](https://www.linkedin.com/in/ethan-gates5329) | [GitHub](#) | [Personal Website](#)

SUMMARY

AI-Native Software Engineer building production AI products, agent systems, and data-intensive full-stack applications. Healthcare SaaS engineer with hands-on work in grounded AI, local inference, orchestration, and evaluation.

EDUCATION

B.S. Computer Science, Western Governors University

Completed: Feb 2025

PROFESSIONAL EXPERIENCE

Software Engineer, DOVAXIS (Healthcare SaaS patient referrals & admissions)

May 2025 – Present

Began as a paid engineering intern; converted to full-time November 2025.

- **Built and operate a multi-provider LLM service (GraphQL on Cloud Run) that internal services call instead of the OpenAI SDK:** a model registry with provider auto-detection across OpenAI and Gemini, a unified quality tier, streaming and tool calling, per-call telemetry for tokens, cost, latency, and organization, plus a gated public chat (pre-launch) with cited web search.
- **Built an evaluation rig for AI clinical-document extraction:** batched runs compared row by row against a ground-truth spreadsheet and the FHIR record, fuzzy matching with reviewer overrides persisted per run, an accuracy dashboard that separates missed data from rows the model invented, a strict mode that penalizes both, and SSN masking during review.
- **Rebuilt how a legacy healthcare app proves changes are safe, with agentic workflows generating and governing the suites:** an offline Firebase emulator with sanitized fixtures so every test runs isolated from live data, browser and Firebase-backed suites wired into a co-designed CI/CD pipeline, and LLM-run audits that turn past incidents into regression checks.
- **Owned a multi-tenant authorization migration end to end:** consolidated scattered role, route, and resource checks into a Permit.io-backed facade with bulk decisions and cached permissions, built the reconciliation report and per-organization readiness verdict that gated the cutover, then ran the production cutover of every customer organization.
- **Built the version and index governance for a multi-tenant clinical document library:** a unified version-history and DocHub lifecycle, re-indexing that keeps the AI search index in step with the library, a per-organization vector audit that finds and clears stale entries, and cited-source attribution so answers separate the documents they used from other hits.

INDEPENDENT PROJECTS

- **GatesAI Chat:** Built and shipped a local-first AI workspace (React/Tauri/Rust): OpenRouter/Ollama routing, multi-round tool execution, sandboxed files/shell/Git/SQLite, memory, and image generation, with automated checks for streams, tools, and user journeys. [\[Repo\]](#)
- **Milo:** Built a local voice companion: whisper.cpp, Ollama, and Pocket TTS in one Python process, interruptible mid-sentence, and an animated robot whose mouth tracks playback. [\[Repo\]](#)
- **Media Vault:** Built a private photo and video library that runs on your own machine: timeline, people, places, trips, and search, with local models for faces, text in photos, image search by description, and video transcripts. [\[Repo\]](#)
- **Fathom:** Built a market-analysis suite in TypeScript and React: portfolio backtesting over 155 years of asset-class history, Monte Carlo retirement simulation, SEC-filing research, and broker-CSV portfolio X-ray. [\[Demo\]](#)

TECHNICAL SKILLS

- **Languages:** TypeScript, Python, Rust, JavaScript, SQL, C#.
- **AI / LLM Engineering:** Agent orchestration, context management, tool calling, grounded retrieval, embeddings and vector search, evaluation and benchmarking, cost and latency telemetry, OpenAI, Gemini, OpenRouter, Ollama.
- **Frameworks & Libraries:** React, MobX, FastAPI, Tauri, TensorFlow, NumPy, Pandas, Tailwind CSS.
- **Tools & Infrastructure:** Git/GitHub, Playwright, Firebase, PostgreSQL, Docker, Vite, Cloud Run, Linux.